Real AI Cases — Paired Analysis Handout

Instructions

Your team receives one pair: one promising implementation and one failure or warning case. Read both and discuss:

1. What value was the organization seeking?
2. What did the AI do, and what did humans do?
3. What evidence supports the account?
4. Which control helped or was missing?
5. Who was accountable for checking the output or system?
6. What should be retained, stopped, or redesigned?

Claims below are concise English summaries of the linked sources. “Promising” does not mean independently proven, risk-free, or transferable to every organization.

Pair A — Grounding and source verification

Promising: FFI and the Norwegian Navy HR assistant

FFI, the Norwegian Navy, and a Norwegian technology company developed a conversational HR assistant. It used retrieval-augmented generation (RAG) to answer questions from a bounded collection of relevant agreements, rules, and procedures. The assistant could show references to this document base. It was first piloted and then tested with line managers in a field experiment. FFI reports that frequent users answered recurring questions faster, especially questions about working time, overtime, travel, leave, and compensation.

This is evidence of usefulness in a defined context, not proof that RAG always produces correct answers. Important features include the narrow purpose, controlled source set, visible references, staged testing, and user feedback.

Source: FFI — field experiment with a digital HR assistant (https://www.ffi.no/publikasjoner/arkiv/okt-produktivitet-med-kunstig-intelligens-et-felteksperiment-med-digital-hr-assistent)

Warning: Asker municipality's legal references

Asker municipality used AI while searching for relevant Supreme Court decisions for a political decision document. VG later reported that one cited decision did not exist and that other decisions did not support the statements attributed to them. The municipality apologized for incorrect references while maintaining that its underlying conclusion was valid.

The case shows that a plausible legal explanation and professional formatting do not verify a source. Even if the final decision might have been unchanged, incorrect references can weaken decision quality, accountability, and public trust. A competent reviewer must open the legal source and verify that it exists and supports the claim.

Source: VG — Asker municipality and incorrect Supreme Court references (https://www.vg.no/nyheter/i/OkyLKA/asker-kommune-brukte-ai-som-diktet-opp-hoeyesterettsdommer)

Pair question: Which verification control was present in the first case but insufficient in the second?

Real AI Cases — Paired Analysis Handout

Instructions

Your team receives one pair: one promising implementation and one failure or warning case. Read both and discuss:

1. What value was the organization seeking?
2. What did the AI do, and what did humans do?
3. What evidence supports the account?
4. Which control helped or was missing?
5. Who was accountable for checking the output or system?
6. What should be retained, stopped, or redesigned?

Claims below are concise English summaries of the linked sources. “Promising” does not mean independently proven, risk-free, or transferable to every organization.

Pair B — Evidence, measurement, and accountability

Promising: Norges Bank Investment Management (NBIM)

NBIM developed an internal AI strategy, provided employees with access to approved leading models, and invested in training through technology days, workshops, and courses. In its 2024 annual report, NBIM states that employees reported an average productivity improvement of 15% from AI tools.

This is relevant organizational evidence, but it is self-reported. It is not the same as a controlled measurement showing that output rose by exactly 15%, nor does it prove the same result for every role or task. A responsible manager preserves that distinction when discussing return on investment.

Source: NBIM — Government Pension Fund Global Annual Report 2024, p. 54 (https://www.nbim.no/contentassets/490f9f062cfc4694b12c45f4d04ab0a5/gpfg_annual_report_2024_uuweb2.pdf)

Failure: Norwegian Police University College

The Norwegian Police University College included Copilot-generated figures in a note comparing how institutions calculated teaching time. Subsequent reporting said that the chat record contained documented figures for only one institution; other figures were based on assumptions or estimates. The rector acknowledged that the information had not been checked properly.

The failure is not only that Copilot generated unsupported numbers. Generated assumptions entered a professional document without being clearly labelled and independently verified. Exact numbers deserve more checking, not more trust, because false precision can appear authoritative.

Source: Digi/NTB — Police University College apologizes for incorrect AI-generated figures (https://www.digi.no/nyhetsstudio/politihoegskolen-beklager-etter-bruk-av-falske-ki-tall/107077?showFeed=1)

Pair question: How can a manager distinguish an observed measure, a self-reported estimate, a reasoned assumption, and a generated guess?

Real AI Cases — Paired Analysis Handout

Instructions

Your team receives one pair: one promising implementation and one failure or warning case. Read both and discuss:

1. What value was the organization seeking?
2. What did the AI do, and what did humans do?
3. What evidence supports the account?
4. Which control helped or was missing?
5. Who was accountable for checking the output or system?
6. What should be retained, stopped, or redesigned?

Claims below are concise English summaries of the linked sources. “Promising” does not mean independently proven, risk-free, or transferable to every organization.

Pair C — Adoption and data governance

Promising: Wallenius Wilhelmsen's Copilot rollout

A Microsoft customer story describes how the Norwegian-headquartered shipping and logistics company approached Copilot adoption. The company conducted a security evaluation, selected participants deliberately, involved senior leaders, and ran a communication and learning campaign. The story reports that usage reached 80% following a seven-week campaign.

This is a vendor-published customer story and does not independently establish productivity or financial return. Its useful lesson is about implementation: access to a tool is not enough. Security evaluation, role modelling, communication, learning, and support influence whether people use a new system responsibly.

Source: Microsoft customer story — Wallenius Wilhelmsen (https://www.microsoft.com/en/customers/story/1786814215980808594-wilh-wilhelmsen-microsoft-viva-goals-travel-and-transportation-en-norway)

Warning: Schibsted integration incident

News reporting stated that a technical integration sent internal comments associated with published Schibsted articles to OpenAI for about a year. Such comments could include editorial notes and source contact information. Schibsted stated that the information had been filtered out, was not exposed to outsiders, and was deleted after the issue was discovered.

The case shows that responsible AI is not only about what an employee types into a chatbot. Data may also flow through APIs, metadata, system configuration, and supplier integrations. Organizations need pre-launch data-flow reviews, access controls, logging, monitoring, and incident procedures.

Source: Dagbladet — internal Schibsted notes included in an OpenAI integration (https://www.dagbladet.no/nyheter/krisemote-delte-interne-notater/84685822)

Pair question: Which technical and organizational checks should occur before and after an AI integration is activated?