

Session 4 Course Note: Trust, Compliance, and Organizational Failure

For students in Innovation and Technology Management

Workshop connection: Session 4, 12:30–13:20

Purpose: Learn to evaluate whether an AI-supported system has earned enough trust for a particular task and consequence.

Legal note: *This is management education, not legal advice. AI regulation is changing. The status summary was checked on 20 July 2026 and should be checked again before being used in professional work.*

How to use this note

Responsible AI is not a separate ethical discussion after the technical design is finished. Trustworthiness, compliance, security, human oversight, and evidence must shape the use case from the beginning.

After reading, you should be able to:

- Explain why model error is only one possible cause of organizational failure.
- Analyse a claim about AI success or failure critically.
- Identify reliability, safety, security, privacy, fairness, transparency, accountability, and sustainability concerns.
- Explain why a nominal “human in the loop” may be ineffective.
- Apply the workshop’s TRUST review.
- Distinguish law, organizational policy, professional duty, and ethical judgment.
- Recommend go, conditional go, or no-go/redesign.

1. Trust must be specific and earned

“Do you trust AI?” is too broad to guide a decision. Ask instead:

Has this system, in this workflow, for this purpose and population, produced enough evidence and control for this level of consequence?

You might reasonably trust one system to suggest alternative wording while refusing to trust the same model to determine access to employment. Trust should be **calibrated**: neither blind confidence nor automatic rejection.

Trustworthy performance depends on more than a model:

- Valid data and sources.
- Testing under realistic conditions.
- Competent users.
- A suitable workflow.
- Effective review and escalation.
- Security and supplier management.
- Transparency to affected people.
- Monitoring after deployment.
- Clear accountability.

NIST describes trustworthy AI through characteristics including validity and reliability, safety, security and resilience, accountability and transparency, explainability and interpretability, privacy enhancement, and fairness with harmful bias managed (NIST AI research (<https://www.nist.gov/ai-research>)). These characteristics

can conflict. For example, collecting more data to monitor fairness may increase privacy risk.

2. From model error to organizational failure

Use this simplified relationship:

Model or data error + weak workflow + unclear ownership = organizational failure

The terms do not always appear together, but the equation directs attention toward management choices.

Model and data contributors

- Incorrect or unsupported generation.
- Poor classification or prediction.
- Biased or unrepresentative training data.
- Outdated knowledge.
- Distribution shift after deployment.
- Adversarial manipulation.

Workflow contributors

- The tool is used for a task it was not evaluated for.
- A draft automatically enters a formal document.
- Review is rushed, optional, or performed without original sources.
- Users cannot report or correct an error.
- The system is integrated with more data or authority than necessary.

Organizational contributors

- No accountable owner.
- Incentives reward speed or adoption but ignore quality.
- Users lack training.
- Procurement relies on vendor claims without testing.
- Responsibility is divided between business, IT, legal, and supplier teams.
- Known incidents do not lead to redesign.

When an AI-generated legal citation enters a decision document, saying “the model hallucinated” describes a technical symptom. It does not explain why the claim was accepted, why the source was not opened, or who approved publication.

3. A practical risk map

Reliability and validity

Reliability asks whether performance is sufficiently consistent. **Validity** asks whether the system measures or produces what the use case actually requires.

A defect classifier may be consistent but invalid if its labels do not represent the engineering decision. A chatbot may answer fluently but fail to retrieve the effective procedure.

Ask:

- What exactly was tested?
- Does the test represent real cases, languages, equipment, and users?
- Which errors are hidden by an average score?
- How will performance change over time?

Safety

Safety concerns physical and psychological harm. Consider incorrect action, delayed action, unavailable systems, unsafe interaction with equipment, and the possibility that humans lose critical skills.

Security and resilience

AI expands the attack surface through data pipelines, external models, APIs, plugins, retrieved documents, and generated output. Prompt injection, data poisoning, insecure output handling, model or supplier compromise, and excessive tool permissions are examples. NIST's GenAI profile and the OWASP GenAI Security Project provide useful starting points (NIST AI 600-1 (<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>); OWASP LLM Top 10 (<https://genai.owasp.org/llm-top-10/>)).

Privacy

Personal data can appear in training data, prompts, retrieved documents, logs, generated output, or monitoring. Privacy questions include purpose, legal basis, data minimization, retention, access, security, individual rights, and international processing.

Datatilsynet recommends risk assessment, privacy by design, regular testing, and explicit support for individual rights when organizations purchase or use AI-based systems (Datatilsynet (<https://www.datatilsynet.no/regelverk-og-verktoy/rapporter-og-utredninger/kunstig-intelligens/anbefalinger/>)).

Fairness and harmful bias

Bias may enter through historical decisions, labels, sampling, proxies, problem definitions, thresholds, language performance, interface design, or unequal access to appeal.

Do not ask only whether two groups receive the same average accuracy. Also ask:

- Which outcome is being distributed?
- Which groups were included in testing?
- Are error types distributed differently?
- Can affected people understand and challenge the result?
- Does the system reproduce an unfair existing process more efficiently?

Transparency and explainability

Different people need different information:

- Users need instructions, limits, and evidence.
- Affected people may need notice, reasons, and a route to challenge.
- Managers need performance, risk, and incident information.
- Auditors or regulators may need logs and documentation.
- Engineers need technical detail sufficient for investigation.

An explanation that sounds understandable may still be inaccurate. Traceability to evidence and decision logic is often more useful than a generated narrative about why the model answered.

Accountability

Accountability requires more than assigning blame after failure. It includes authority and duties before deployment:

- Who approves the purpose and scope?
- Who owns data and source quality?
- Who accepts residual risk?
- Who monitors performance?
- Who handles incidents and complaints?

- Who can pause or stop the system?

Sustainability

AI has environmental and resource costs involving computing, energy, water, hardware, and supply chains. The relevant management question is proportionality: is the resource-intensive solution justified by the value, and can a smaller model or simpler method perform the task? The 2026 Stanford AI Index documents growing infrastructure and environmental demands alongside technical progress (Stanford HAI, 2026—Research and Development (<https://hai.stanford.edu/ai-index/2026-ai-index-report/research-and-development>)).

4. Human oversight is a designed capability

“A human makes the final decision” can provide false confidence.

Automation bias occurs when people defer excessively to automated advice and reduce independent information seeking. A systematic review found that decision support can introduce new errors and that task complexity, workload, and verification difficulty affect over-reliance (Goddard, Roudsari, and Wyatt (<https://pmc.ncbi.nlm.nih.gov/articles/PMC7651899/>)). NIST also identifies human–AI configuration and over-reliance as GenAI risk areas.

Effective oversight requires:

- **Competence:** The reviewer understands the domain and the system’s limits.
- **Evidence:** The reviewer can inspect the original source or data.
- **Time:** Review is realistic within workload and deadlines.
- **Authority:** The reviewer can reject, override, or escalate.
- **Independence:** The review is not merely asking the same model to approve itself.
- **Feedback:** Corrections reach the people responsible for the process.
- **Accountability:** A named role owns the final action.

Interface design matters. If the AI recommendation appears first and dominates attention, the human may anchor on it. For important decisions, it may be better for the human to record an initial assessment before seeing the AI recommendation.

5. How to read an AI success claim

Ask six questions:

1. **Who published the claim?** A vendor, customer, independent researcher, regulator, or journalist?
2. **What was measured?** Access, usage, speed, quality, outcome, financial value, or satisfaction?
3. **How was it measured?** Controlled experiment, operational comparison, self-report, testimonial, or estimate?
4. **What is the comparison?** Previous performance, a control group, another tool, or no comparison?
5. **What context is missing?** Task, users, language, time, costs, harms, and failed cases?
6. **Who was accountable for checking the output or system?**

Example: NBIM’s productivity statement

Norges Bank Investment Management’s 2024 annual report says that employees reported an average 15% productivity improvement from AI tools. The report also describes access to models, an internal strategy, learning activities, and technology days (NBIM Annual Report 2024, p. 54 (https://www.nbim.no/contentassets/490f9f062cfc4694b12c45f4d04ab0a5/gpfg_annual_report_2024_uuweb2.pdf)).

This is relevant evidence about employee perceptions and organizational activity. It is not a controlled estimate of output, quality, profit, or causality. The correct response is not to dismiss the number, but to label the evidence accurately.

6. How to read an AI failure story

News reports are useful for identifying warning patterns, but they may not contain complete technical evidence. Separate:

- Confirmed events.
- Claims by involved parties.
- Journalistic findings.
- Your interpretation.
- Information not publicly available.

Avoid using failure cases as entertainment. Ask what management could have changed before the event and what learning should follow.

Asker municipality: grounding and source checking

Reporting found that incorrect or non-existent Supreme Court references appeared in political decision material after AI-assisted legal searching (VG (<https://www.vg.no/nyheter/i/OkyLKA/asker-kommune-brukte-ai-som-diktet-opp-hoeyesterettsdommer>)). The management issue is not merely generation. It includes use of authoritative legal databases, citation verification, reviewer competence, and document approval.

Norwegian Police University College: evidence categories

Copilot-generated figures appeared in a note comparing teaching time, while reporting later found documented data for only one institution. The organization acknowledged inadequate verification (Digi/NTB (<https://www.digi.no/nyhetsstudio/politihoeegskolen-beklager-etter-bruk-av-falske-ki-tall/107077?showFeed=1>)). The case demonstrates the importance of separating observed measures, assumptions, estimates, and generated guesses.

Schibsted: integration and data flow

A technical integration reportedly sent internal comments connected to published articles to OpenAI over an extended period. The company stated that the data were filtered, not exposed to external users, and deleted (Dagbladet (<https://www.dagbladet.no/nyheter/krisemote-delte-interne-notater/84685822>)). The case directs attention toward APIs, metadata, supplier configurations, logs, and continuous monitoring—not only employee prompts.

7. Compliance: four layers to distinguish

Law

Binding requirements may arise from AI regulation, privacy, employment, equality, product safety, cybersecurity, intellectual property, contract, administrative law, or sector-specific rules.

Organizational policy

An employer or university may set stricter rules than the general law, including approved tools, data classifications, review, or disclosure.

Professional and contractual duty

Engineers, managers, lawyers, auditors, researchers, health professionals, and others may have duties defined by professional standards, contracts, or role-specific expectations.

Ethics and legitimacy

An action may be legal yet still unfair, manipulative, unsafe, environmentally disproportionate, or inconsistent with organizational values and public trust.

Compliance is the minimum condition, not the complete definition of responsible use.

8. The EU AI Act and Norway: status checked 20 July 2026

The EU AI Act uses a risk-based structure. It includes prohibited practices, obligations for general-purpose AI, transparency requirements, and more extensive requirements for certain high-risk uses.

According to the European Commission's current implementation page:

- The Act entered into force in 2024.
- Prohibited-practice and AI-literacy provisions began applying in February 2025.
- Governance and general-purpose-model obligations began applying in August 2025.
- Many provisions and transparency obligations apply from 2 August 2026.
- Following the May 2026 political agreement on simplification, high-risk-system timelines are later and depend on category.

Use the European Commission's live AI Act page (<https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>) rather than copying an old timeline. The Commission published additional transparency guidance on 20 July 2026, shortly before the workshop (Article 50 transparency guidance announcement (<https://digital-strategy.ec.europa.eu/en/news/commission-publishes-guidelines-transparency-obligations-providers-and-deployers-certain-ai-systems>)).

Norway is not an EU member and implementation requires EEA incorporation and Norwegian legislation. A Norwegian government update in May 2026 stated that high-risk rules would apply in Norway only after an AI law had been adopted by the Storting; the government was preparing a legislative proposal and intended implementation after EEA incorporation (Norwegian government, 26 May 2026 (<https://www.regjeringen.no/no/aktuelt/send-innspeil-til-kva-eit-hogrisiko-ki-system-skal-verre-innan-23.-juni/id3161303/>)). The earlier Norwegian consultation remains listed as under consideration (consultation page (<https://www.regjeringen.no/no/dokumenter/3112327/id3112327/>)).

For management purposes, do not wait for a legal deadline before establishing data governance, competence, documentation, testing, human oversight, and incident procedures.

9. USN responsibilities

USN's AI policy and teaching guidance emphasize transparency, quality assurance, protection of personal and confidential information, attention to bias and sustainability, and retained human responsibility. For assessed work, the rules of the specific course and assessment apply. Check them rather than assuming that all AI use is either permitted or prohibited (USN guidance (<https://edu.usn.no/utvikle-din-undervisning/kunstig-intelligens/laering-undervisning-og-vurdering-med-ki/>)).

10. Apply the TRUST review

The workshop uses five connected areas.

T — Task and consequence

- Is AI suitable for this task?
- What happens if it is wrong, unavailable, delayed, manipulated, or incomplete?
- Is the result reversible?
- Is the autonomy proportionate?

R — Rights and rules

- Could the use affect privacy, employment, equality, safety, intellectual property, contracts, or access to services?
- Which legal, policy, professional, and ethical duties apply?
- Can affected people obtain information and challenge an outcome?

U — Users and oversight

- Who uses the system and who is affected by it?
- Who checks, approves, overrides, and escalates?
- Do they have competence, evidence, time, and authority?
- How will automation bias and deskilling be reduced?

S — Sources, security, and sustainability

- Are sources permitted, current, representative, and traceable?
- Are access controls preserved?
- Could input manipulate the system?
- Are supplier, resilience, and environmental costs proportionate?

T — Testing and transparency

- Which value, quality, fairness, and risk measures are tested?
- Which versions, inputs, outputs, decisions, and incidents are logged?
- Who must be told that AI was used?
- How are drift, misuse, and complaints monitored?

Use the TRUST worksheet (in the shared workshop materials) to identify the most important concern and a concrete safeguard for each area.

11. Pre-mortem and decision

A pre-mortem assumes that the initiative has failed and asks what happened:

Six months after launch, the initiative caused an operational, legal, ethical, security, or reputational failure. What is the most credible explanation?

This makes it easier to discuss failure without requiring anyone to attack the current proposal personally.

For the three most credible scenarios, define:

- Trigger or contributing condition.
- Harm and affected stakeholder.
- Detection method.
- Preventive safeguard.
- Response owner.

Then decide:

- **Go:** sufficiently bounded for a limited test.
- **Conditional go:** proceed only after named conditions are met.
- **No-go/redesign:** task, data, autonomy, control, or consequence is unsuitable.

An ethical or well-governed project is not automatically effective. A responsible decision considers both value and harm.

12. Governance across the lifecycle

Before deployment

- Approve the purpose and scope.
- Map data, suppliers, and affected stakeholders.

- Define measures, roles, exclusions, and stop conditions.
- Test normal, edge, and adversarial cases.
- Prepare communication, training, and incident procedures.

During use

- Monitor quality, risk, usage, overrides, and complaints.
- Review changes to models, data, prompts, integrations, and suppliers.
- Ensure human checks actually occur.
- Investigate differences across contexts and groups.

After incidents or retirement

- Contain harm.
- Preserve evidence and logs.
- Inform responsible roles and affected parties as required.
- Correct downstream records.
- Learn and redesign.
- Remove access, integrations, and retained data when the system is retired.

Governance is an operating process, not a document produced for approval day.

13. Reflection questions

1. When is a citation sufficient evidence, and when must you inspect the source and method behind it?
2. Which condition—competence, evidence, time, authority, independence, feedback, or accountability—is most often missing from “human oversight”?
3. Select one workshop case. What was the model’s role, the human role, the missing or effective control, and the accountable owner?
4. Name one use that could be legally permitted but ethically unacceptable.
5. What evidence would cause you to change a go decision to conditional go or no-go?

Key terms

- **Trustworthiness:** Demonstrated qualities that justify reliance for a defined context and consequence.
- **Automation bias:** Excessive human deference to automated advice.
- **Drift:** Change in data, context, or performance after deployment.
- **Accountability:** Assigned authority and obligation to explain, monitor, correct, and accept responsibility.
- **Transparency:** Making relevant information about AI use, limits, evidence, and decisions available to the appropriate audience.
- **Explainability:** The extent to which reasons or mechanisms behind an output can be understood.
- **Pre-mortem:** A structured exercise that assumes failure and works backward to identify causes and safeguards.
- **Residual risk:** Risk remaining after safeguards are applied.

References and further reading

- NIST. AI Risk Management Framework (<https://www.nist.gov/itl/ai-risk-management-framework>) and Generative AI Profile (<https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence>).

- European Commission. AI Act implementation page (<https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>).
- Norwegian Government. Consultation on a Norwegian AI law (<https://www.regjeringen.no/no/dokumenter/3112327/id3112327/>) and May 2026 implementation update (<https://www.regjeringen.no/no/aktuelt/send-innspeil-til-kva-eit-hogrisiko-ki-system-skall-verre-innan-23-juni/id3161303/>).
- Datatilsynet. Recommendations for AI-based systems (<https://www.datatilsynet.no/regelverk-og-verktoy/rapporter-og-utredninger/kunstig-intelligens/anbefalinger/>).
- Goddard, K., Roudsari, A., and Wyatt, J. Automation bias and verification complexity: a systematic review (<https://pmc.ncbi.nlm.nih.gov/articles/PMC7651899/>).
- Stanford HAI. Responsible AI chapter, 2026 AI Index (<https://hai.stanford.edu/ai-index/2026-ai-index-report/responsible-ai>).
- USN. Guidance on AI in teaching, learning, and assessment (<https://edu.usn.no/utvikle-din-undervisning/kunstig-intelligens/laering-undervisning-og-vurdering-med-ki/>).

Session takeaway

Trust is not a general attitude toward AI. It is a conclusion supported by evidence, competent people, effective controls, and accountable ownership for a defined use.